

Genetic Programming

Rui Mendes

Advantages of Genetic Programming

- Allows the automatic creation of programs
- Evolves programs that generate solutions instead of the solutions themselves
- More generic
- More powerful

Disadvantages of Genetic Programming

- Needs more CPU time for training
 - ▶ Evaluating programs typically takes longer than evaluating solutions
 - ▶ Programs need to be tested in as many situations as possible
- It is often hard to understand the programs that were generated
- Sometimes, it is necessary to use simplifiers in order to be able to understand the result

Representations

- Trees** A tree represents a program where branches represent non-terminals and leaves represent terminals. The number of sub-trees of a branch is the arity of the function
- Linear** Similar to machine code. Each instruction uses registers for its arguments and storing the result
- Graphs** The program is represented by a graph. Data is usually passed using the stack. The execution flow uses the edges of the graph

Advantages of Tree Representation

- Intuitive structure for representing programs
- Very similar to Lisp or Syntax Trees
- Very easy to manipulate by genetic operators
- Easy to visualize

Terminals

- Functions of zero arity
- They represent:
 - 1 The program inputs
 - 2 The constants (Random Ephemeral Constants) that are numbers that are generated when GP is initialized and used by all the individuals
 - 3 Functions with no parameters but that produce side effects

Functions

Booleans And, Or, Not, Xor

Arithmetic $+$, $-$, \times , $/$

Transcending $\sin x$, $\cos x$, $\log x$, e^x

Attribution (Assign a 1), (Read a)

Conditionals If Then Else

Composing functions (prog2 fun1 fun2)

Others read sensor, turn left, turn right, move ahead

Initialization

Full Only non-terminals are used until the maximum depth of the tree is reached. At that time, only terminals are used.

Grow When generating a node, both terminals and non-terminals are used. This generates irregular trees


Ramped half and half Method that combines **Full** and **Grow**

Ramped half and half


- 1 The same number of trees is generated for all depths ranging from 2 until the maximum depth
- 2 For each depth, half of the trees are generated using the *Full* technique while the other half uses the *Grow* technique
- 3 The same number of trees is generated for all depths ranging from 2 until the maximum depth

Initialization types

Full



Grow




Crossover


- Two parents are selected for exchanging genetic material
- A subtree is selected at random in each parent
- The subtrees are exchanged between parents
- The exchange is only performed if it doesn't violate the maximum depth constraint
- One of the advantages of crossover between trees is that it can generate different offspring even when both parents have the same information

Crossover


Parent 1




Parent 2



Child 1




Child 2




Crossover between equal individuals


Parent 1




Parent 2



Child 1



Child 2




Mutation


- One individual is chosen to undergo mutation
- A subtree is selected at random
- The subtree is exchanged by a random subtree
- The random subtree is generated while enforcing the maximum depth constraint
- Another possible mutation is to substitute a node by another of the same arity

Mutation of a subtree

Before




After




Mutation of a node

Before



After



Fitness

Fitness A measure of the degree of success of a program

Continuity Small improvements in the program should produce a slightly better fitness; large improvements should translate into a much larger fitness

Standardized fitness Where 0 is given to the optimal solution

Normalized fitness Where the fitness is always between 0 and 1

Examples of Fitness

- The number of collisions of a robot learning how to avoid obstacles
- The average speed of a robot who should learn how to avoid obstacles and always keep in motion
- The error in classification or regression tasks
- The money won by an agent who is learning how to play Poker
- The amount of food gathered by an agent in an artificial life scenario
- The number of goals scored by a robot playing soccer
- The number of hits scored by an agent in a game
- The number of turns an agent survives in a game

Selection

Proportional to the fitness $p_i = \frac{f_i}{\sum_i f_i}$

(μ, λ) selection λ individuals are generated and the best μ ones are selected

Tournament Several individuals are randomly selected from the population and the best ones create offspring by using the genetic operators and substitute the worse ones

- Two individuals are randomly chosen and the best one creates a mutated copy that substitutes the worse one
- Four individuals are randomly chosen and the best two undergo *crossover* and *mutation* to generate two offspring that substitute the worse ones

Generational GP

- ① Initialize the population
- ② Evaluate all the individuals and compute their fitness
- ③ Until the new population is complete:
 - ① Select two parents
 - ② Apply genetic operators
 - ③ Insert offspring into the new population
- ④ Repeat steps 2 and 3 until reaching the termination criterion
- ⑤ Present the best individual in the population

Steady-State GP

- ➊ Initialize the population
- ➋ Select a random subset of the population to undergo the tournament
- ➌ Rank the individuals
- ➍ Apply the genetic operators to the winners of the tournament
- ➎ Substitute the losers by the offspring
- ➏ Repeat the previous steps until reaching the stopping criterion
- ➐ Present the best individual in the population

Advantages of SSGP

- It is not necessary to evaluate all the individuals
- It is only necessary to rank all individuals in the tournament
- This can be done by a competition between the individuals without having a numeric value for the fitness
- It is possible to divide the population in non communicating islands thus parallelizing the algorithm

Island Model

- The population is divided into several islands
- Each island is standalone
- Some individuals migrate from time to time between islands
- The most usual model uses a circular pattern
- This model often has better results than using the panmictic model

Parsimony

- Functions used by GP should be sufficiently powerful to encode problem solutions
- Care should be taken in order to avoid using more functions than is strictly necessary
- This increases the problem space and hardens the task of searching for solutions

Function Closure

- Functions should be able to work with any arguments
- Division must be protected
- This type of protection must be performed for all functions (e.g., square roots or logarithms)

Parameters

Configuration Parameters

- Population Size
- Initialization Type
- Genetic Operators
- Operator Probabilities

Maximum Size

- A big program takes a long time to run
- A big program can encode very specific solutions
- Very specific solutions are hard to generalize
- Maximum number of nodes
- Maximum Depth
- Maximum Number of Instructions

Applications

- Curve fitting, data modelling
- Image and signal processing
- Time series prediction
- Industrial process control
- Ant foraging in artificial life environments
- Chess endgame players
- Evolve Robocode players
- Heuristics for a Monte Carlo Tree Search Ms Pac-Man agent
- Automatic StarCraft Strategy Generation